

```
!pip install -q kaggle
!pip install -q tensorflow matplotlib seaborn scikit-learn opencv-python
```

```
# Create Kaggle directory
!mkdir -p ~/.kaggle

# Move kaggle.json if it is already uploaded
!cp kaggle.json ~/.kaggle/

# Set correct permissions
!chmod 600 ~/.kaggle/kaggle.json

# Verify
!ls ~/.kaggle

kaggle.json
```

```
!kaggle datasets download -d tawsifurrahman/covid19-radiography-database
```

```
Dataset URL: https://www.kaggle.com/datasets/tawsifurrahman/covid19-radiography-database
License(s): copyright-authors
Downloading covid19-radiography-database.zip to /content
96% 746M/778M [00:03<00:00, 208MB/s]
100% 778M/778M [00:03<00:00, 210MB/s]
```

```
!unzip -q covid19-radiography-database.zip
!ls
```

```
covid19-radiography-database.zip  kaggle.json
COVID-19_Radiography_Dataset      sample_data
```

```
import os
import cv2
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt
import seaborn as sns
import tensorflow as tf
from tensorflow.keras import layers, Model
from tensorflow.keras.applications import EfficientNetB3
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix

IMG_SIZE = 224
BATCH_SIZE = 16
EPOCHS = 15

CLASSES = ['COVID']

print("TensorFlow:", tf.__version__)
print("GPU:", tf.config.list_physical_devices('GPU'))
```

```
TensorFlow: 2.19.0
GPU: [PhysicalDevice(name='/physical_device:GPU:0', device_type='GPU')]
```

```
DATASET_PATH = "COVID-19_Radiography_Dataset"
```

```
def load_images(base_path):
    X, y = [], []

    for idx, cls in enumerate(CLASSES):
        img_dir = os.path.join(base_path, cls, "images")
        print(f"Loading images from class: {cls}")

        # Safety check
        if not os.path.exists(img_dir):
            print(f"⚠️ Directory not found: {img_dir}")
            continue

        for img_name in os.listdir(img_dir):
            img_path = os.path.join(img_dir, img_name)

            # Read image
            img = cv2.imread(img_path)
            if img is None:
                continue # skip corrupted or non-image files
```

```
# Preprocessing
img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
img = cv2.resize(img, (IMG_SIZE, IMG_SIZE))
img = img.astype(np.float32) / 255.0

X.append(img)
y.append(idx)

return np.array(X, dtype=np.float32), np.array(y)
```

```
# Load dataset
X, y = load_images(DATASET_PATH)
print("✅ Dataset loaded successfully!")
print("X shape:", X.shape)
print("y shape:", y.shape)
```

Loading images from class: COVID  
 ✅ Dataset loaded successfully!  
 X shape: (3616, 224, 224, 3)  
 y shape: (3616, )

```
X_train, X_temp, y_train, y_temp = train_test_split(
    X, y, test_size=0.3, stratify=y, random_state=42)

X_val, X_test, y_val, y_test = train_test_split(
    X_temp, y_temp, test_size=0.5, stratify=y_temp, random_state=42)

y_train_cat = tf.keras.utils.to_categorical(y_train, 4)
y_val_cat = tf.keras.utils.to_categorical(y_val, 4)
y_test_cat = tf.keras.utils.to_categorical(y_test, 4)

print("Train:", X_train.shape, "Val:", X_val.shape, "Test:", X_test.shape)
```

Train: (2531, 224, 224, 3) Val: (542, 224, 224, 3) Test: (543, 224, 224, 3)

```
def transformer_encoder(x, num_heads=8, key_dim=64, ff_dim=512, dropout=0.2):
    attn = layers.MultiHeadAttention(
        num_heads=num_heads,
        key_dim=key_dim,
        dropout=dropout
    )(x, x)

    x = layers.Add()([x, attn])
    x = layers.LayerNormalization()(x)

    ff = layers.Dense(ff_dim, activation='relu')(x)
    ff = layers.Dense(x.shape[-1])(ff)

    x = layers.Add()([x, ff])
    x = layers.LayerNormalization()(x)

    return x
```

```
def build_model():
    inputs = layers.Input(shape=(224, 224, 3))

    # CNN FEATURE EXTRACTOR
    base = EfficientNetB3(
        include_top=False,
        weights='imagenet',
        input_tensor=inputs
    )

    for layer in base.layers[:-30]:
        layer.trainable = False

    x = base.output
    x = layers.GlobalAveragePooling2D()(x) # 1536-dim

    # ATTENTION MODULE
    q = layers.Dense(512)(x)
    k = layers.Dense(512)(x)
    v = layers.Dense(512)(x)

    q = layers.Reshape((1, 512))(q)
    k = layers.Reshape((1, 512))(k)
    v = layers.Reshape((1, 512))(v)

    x = layers.MultiHeadAttention(
        num_heads=8, key_dim=64
    )(q, k, v)
```

```
j(4, k, v)
# TRANSFORMER ENCODER STACK (4 LAYERS)
for _ in range(4):
    x = transformer_encoder(x)

x = layers.GlobalAveragePooling1D()(x)

# CLASSIFICATION HEAD
x = layers.Dense(256, activation='relu')(x)
x = layers.Dropout(0.3)(x)
x = layers.Dense(128, activation='relu')(x)
x = layers.Dropout(0.3)(x)

outputs = layers.Dense(4, activation='softmax')(x)

return Model(inputs, outputs, name="Hybrid_CNN_Transformer")

model = build_model()
model.compile(
    optimizer=tf.keras.optimizers.Adam(1e-4),
    loss='categorical_crossentropy',
    metrics=['accuracy']
)
model.summary()
```

**Show hidden output**

```
history = model.fit(
    X_train, y_train_cat,
    validation_data=(X_val, y_val_cat),
    epochs=EPOCHS,
    batch_size=BATCH_SIZE
)

Epoch 1/15
/usr/local/lib/python3.12/dist-packages/keras/src/ops/nn.py:944: UserWarning: You are using a softmax over axis 3 of a tensor
warnings.warn(
159/159 ━━━━━━━━━━ 152s 485ms/step - accuracy: 0.9630 - loss: 0.1294 - val_accuracy: 1.0000 - val_loss: 0.0138
Epoch 2/15
159/159 ━━━━━━ 10s 62ms/step - accuracy: 1.0000 - loss: 1.6876e-04 - val_accuracy: 1.0000 - val_loss: 9.2591e-05
Epoch 3/15
159/159 ━━━━━━ 10s 61ms/step - accuracy: 1.0000 - loss: 4.4001e-05 - val_accuracy: 1.0000 - val_loss: 3.2112e-05
Epoch 4/15
159/159 ━━━━━━ 10s 62ms/step - accuracy: 1.0000 - loss: 2.5542e-05 - val_accuracy: 1.0000 - val_loss: 0.0000e+00
Epoch 5/15
159/159 ━━━━━━ 10s 61ms/step - accuracy: 1.0000 - loss: 1.3488e-05 - val_accuracy: 1.0000 - val_loss: 0.0000e+00
Epoch 6/15
159/159 ━━━━━━ 10s 63ms/step - accuracy: 1.0000 - loss: 1.1949e-05 - val_accuracy: 1.0000 - val_loss: 0.0000e+00
Epoch 7/15
159/159 ━━━━━━ 10s 63ms/step - accuracy: 1.0000 - loss: 5.2605e-06 - val_accuracy: 1.0000 - val_loss: 0.0000e+00
Epoch 8/15
159/159 ━━━━━━ 10s 63ms/step - accuracy: 1.0000 - loss: 7.4016e-06 - val_accuracy: 1.0000 - val_loss: 0.0000e+00
Epoch 9/15
159/159 ━━━━━━ 10s 63ms/step - accuracy: 1.0000 - loss: 6.8786e-06 - val_accuracy: 1.0000 - val_loss: 0.0000e+00
Epoch 10/15
159/159 ━━━━━━ 10s 63ms/step - accuracy: 1.0000 - loss: 4.1560e-06 - val_accuracy: 1.0000 - val_loss: 0.0000e+00
Epoch 11/15
159/159 ━━━━━━ 10s 61ms/step - accuracy: 1.0000 - loss: 3.8170e-06 - val_accuracy: 1.0000 - val_loss: 0.0000e+00
Epoch 12/15
159/159 ━━━━━━ 10s 63ms/step - accuracy: 1.0000 - loss: 3.8004e-06 - val_accuracy: 1.0000 - val_loss: 0.0000e+00
Epoch 13/15
159/159 ━━━━━━ 10s 63ms/step - accuracy: 1.0000 - loss: 2.2768e-06 - val_accuracy: 1.0000 - val_loss: 0.0000e+00
Epoch 14/15
159/159 ━━━━━━ 10s 63ms/step - accuracy: 1.0000 - loss: 2.7548e-06 - val_accuracy: 1.0000 - val_loss: 0.0000e+00
Epoch 15/15
159/159 ━━━━━━ 10s 64ms/step - accuracy: 1.0000 - loss: 2.6851e-06 - val_accuracy: 1.0000 - val_loss: 0.0000e+00
```

```
y_pred = np.argmax(model.predict(X_test), axis=1)

print(classification_report(y_test, y_pred, target_names=CLASSES))
```

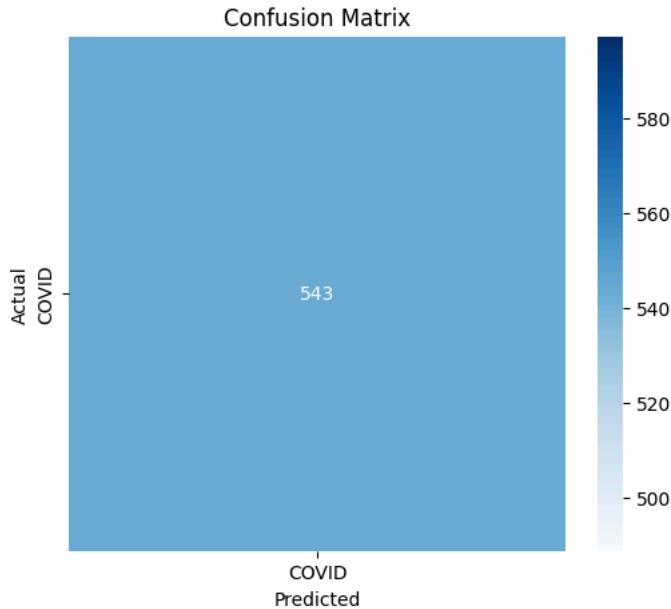
```
/usr/local/lib/python3.12/dist-packages/keras/src/ops/nn.py:944: UserWarning: You are using a softmax over axis 3 of a tensor
warnings.warn(
17/17 ━━━━━━ 53s 1s/step
      precision    recall   f1-score   support
      COVID       1.00      1.00      1.00      543
      accuracy          1.00          1.00          1.00      543
      macro avg       1.00      1.00      1.00      543
```

weighted avg	1.00	1.00	1.00	543
--------------	------	------	------	-----

```
cm = confusion_matrix(y_test, y_pred)
```

```
plt.figure(figsize=(6,5))
sns.heatmap(
    cm, annot=True, fmt='d',
    xticklabels=CLASSES,
    yticklabels=CLASSES,
    cmap='Blues'
)
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```

```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:407: UserWarning: A single label was found in 'y_
warnings.warn(
```



```
model.save("Hybrid_CNN_Transformer_COVID19.h5")
print("Model saved successfully!")
```

```
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file for
Model saved successfully!
```